

Basic Path Finding

Path Planning

*Why not pre-compute
all paths?*


- Path Planning
 - needs to be very fast (especially for games with many characters)
 - needs to generate believable paths



Path Planning

- Path Planning in
 - partially-known environments is a repeated process
 - dynamic environments is also a repeated process

re-planning path as map becomes known





Definition of Path Planning

- Task:
find a feasible (and cost-minimal) path from the current pose to a goal pose
 - Two types of constraints:
environmental constraints (e.g., obstacles)
dynamics/kinematics constraints
 - Generated motion/path should (objective):
be a feasible path
minimize cost such as distance, time, unrealistic effects, ...
-


Path Planning

Examples (of what is usually referred to as path planning):



Path Planning

Examples (of what is usually referred to as motion planning):




Piano Movers' problem

the example above is borrowed from www.cs.cmu.edu/~awm/tutorials

Configuration Space


- **Configuration is legal** if it does not intersect any obstacles and is valid (e.g., does not intersect itself, joint angles are within their limits)
- **Configuration Space** is the set of legal configurations



Configuration Space


- **Configuration is legal** if it does not intersect any obstacles and is valid (e.g., does not intersect itself, joint angles are within their limits)
- **Configuration Space** is the set of legal configurations

What is the dimensionality of this configuration space?



C-Space Transform


- Configuration space for rigid-body objects in 2D world is:
 - 2D if object is circular




- expand all obstacles by the radius of the object r
- planning can be done for a point R (and not a circle anymore)

C-Space Transform

- Configuration space for rigid-body objects in 2D world is:
 - 2D if object is circular



$\xrightarrow{\text{C-Space Transform}}$




Is this a correct expansion?

- expand all obstacles by the radius of the object r
- planning can be done for a point R (and not a circle anymore)

C-Space Transform

- Configuration space for rigid-body objects in 2D world is:
 - 2D if object is circular




- advantage: planning is faster for a single point
- disadvantage: need to expand obstacles every time map is updated (O(n) methods exist to compute distance transforms)

why?

C-Space Transform

- Configuration space for arbitrary objects in 2D world is:
 - 3D if object is non-circular



- advantage: planning is faster for a single point
- disadvantage: constructing C-space is expensive

Planning as Graph Search Problem

1. Construct a graph representing the planning problem
2. Search the graph for a (hopefully, close-to-optimal) path

The two steps above are often interleaved

Planning as Graph Search Problem

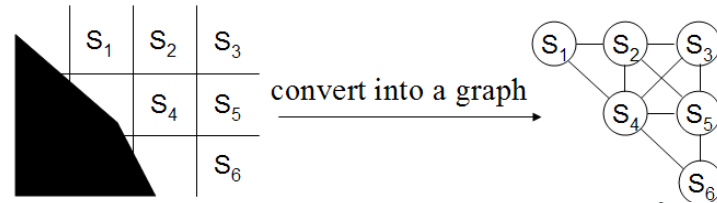
1. Construct a graph representing the planning problem
2. Search the graph for a (hopefully, close-to-optimal) path

The two steps above are often interleaved


Graph Construction

- Cell decomposition

- X-connected grids



- lattice-based graphs



- Skeletonization of the environment/C-Space

- Visibility graphs


- Voronoi diagrams

- Probabilistic roadmaps

- Navmeshes


Graphs Construction

- Once a graph is constructed, we will search it for a least-cost path
- Once again: depending on the planning algorithm, graph construction can be interleaved with graph search




Planning via Cell Decomposition

- Exact Cell Decomposition:
 - overlay convex exact polygons over the free C-space
 - construct a graph, search the graph for a path
 - overly expensive for non-trivial environments and/or above 2D




Planning via Cell Decomposition

- Approximate Cell Decomposition:
 - overlay uniform grid over the C-space (discretize)




Planning via Cell Decomposition

- Approximate Cell Decomposition:
 - construct a graph and search it for a least-cost path




Planning via Cell Decomposition

- Approximate Cell Decomposition:
 - construct a graph and search it for a least-cost path




Planning via Cell Decomposition

- Approximate Cell Decomposition:
 - construct a graph and search it for a least-cost path
 - VERY popular due to its simplicity
 - expensive in high-dimensional spaces
- construct the grid on-the-fly, i.e. while planning – still expensive




Planning via Cell Decomposition

- Approximate Cell Decomposition:
 - what to do with partially blocked cells?



Planning via Cell Decomposition


- Approximate Cell Decomposition:
 - what to do with partially blocked cells?
 - make it untraversable – incomplete (may not find a path that exists)




Planning via Cell Decomposition

- Approximate Cell Decomposition:
 - what to do with partially blocked cells?
 - make it traversable – unsound (may return invalid path)

so, what's the solution?




convert into a graph



search the graph
for a least-cost path
from s_{start} to s_{goal}

Planning via Cell Decomposition


- Approximate Cell Decomposition:
 - solution 1:
 - make the discretization very fine
 - expensive, especially in high-D




Planning via Cell Decomposition

- Approximate Cell Decomposition:
 - solution 2:
 - make the discretization adaptive
 - various ways possible

How?



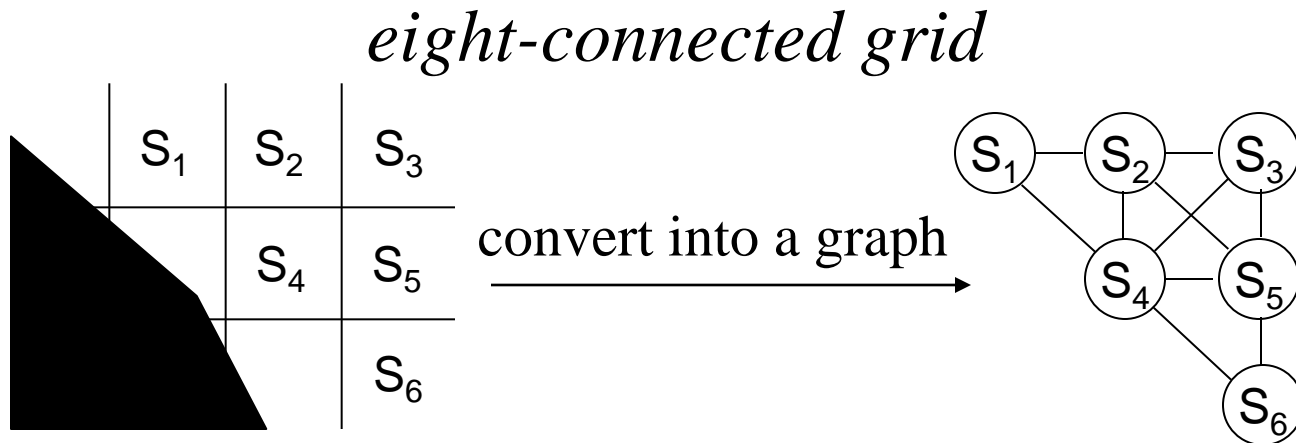
convert into a graph



search the graph
for a least-cost path
from s_{start} to s_{goal}


Planning via Cell Decomposition

- Graph construction:
 - connect neighbors



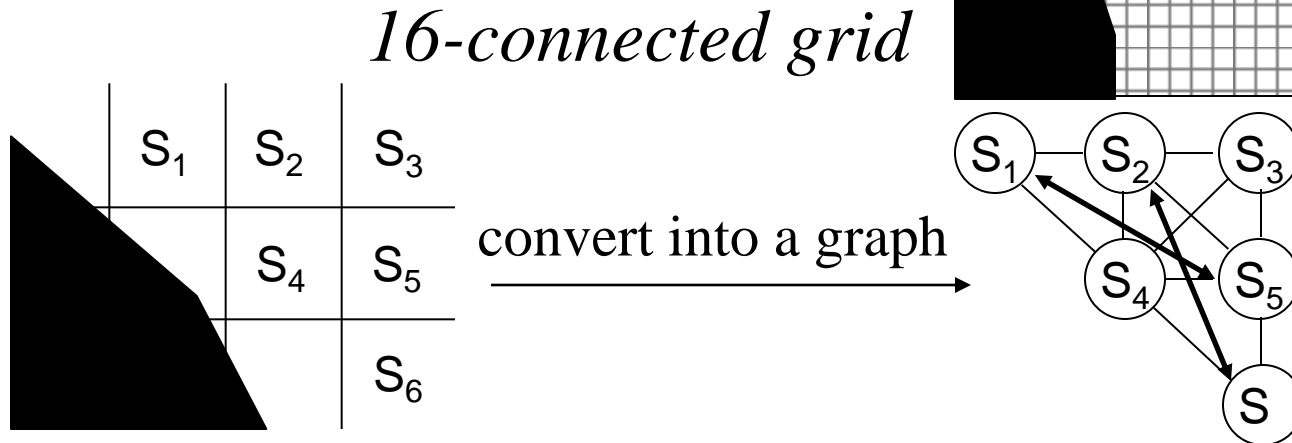
Planning via Cell Decomposition

- Graph construction:
 - connect neighbors
 - path is restricted to 45° degrees



Planning via Cell Decomposition


- Graph construction:
 - connect cells to neighbor of neighbors
 - path is restricted to 22.5° degrees




Planning via Cell Decomposition

- Graph construction:
 - connect cells to neighbor of neighbors
 - path is restricted to 22.5° degrees


Disadvantages?



16-connected grid



convert into a graph



Planning via Cell Decomposition

- Graph construction:
 - lattice graph for computing smooth (realistic) paths


outcome state is the center of the corresponding cell

*each transition is feasible
(constructed beforehand)*

action template




*replicate it
online*




Planning via Cell Decomposition

- Graph construction:
 - lattice graph
 - pros: sparse graph, feasible paths
 - cons: possible incompleteness

action template




*replicate it
online*



Planning via Cell Decomposition

- Graph construction:
 - lattice graph




Planning via Cell Decomposition

- Graph construction:
 - lattice graph

planning on 4D lattice graph:

each state represents $\langle x, y, \text{orientation}, \text{velocity} \rangle$

each edge represents a short feasible motion between corresponding cells




Skeletonization of the C-Space

Skeletonization: construction of a unidimensional representation of the C-space

- Visibility graph
- Voronoi diagram
- Probabilistic road-map
- Navmeshes

Planning via Skeletonization

- **Visibility Graphs** [Wesley & Lozano-Perez '79]
 - based on idea that the shortest path consists of obstacle-free straight line segments connecting all obstacle vertices and start and goal




Planning via Skeletonization

- Visibility Graphs

- based on idea that the shortest path consists of obstacle-free straight line segments connecting all obstacle vertices and start and goal

Assumption?




Planning via Skeletonization

- Visibility Graphs


- based on idea that the shortest path consists of obstacle-free straight line segments connecting all obstacle vertices and start and goal

Assumption?



Planning via Skeletonization

- **Visibility Graphs** [Wesley & Lozano-Perez '79]
 - construct a graph by connecting all vertices, start and goal by obstacle-free straight line segments (graph is $O(n^2)$, where n - # of vert.)
 - search the graph for a shortest path




Planning via Skeletonization

- Visibility Graphs

- advantages:
 - independent of the size of the environment
- disadvantages:
 - path is too close to obstacles
 - hard to deal with non-uniform cost function
 - hard to deal with non-polygonal obstacles

Planning via Skeletonization

- Voronoi diagrams [Rowat '79]
 - voronoi diagram: set of all points that are equidistant to two nearest obstacles
 - based on the idea of maximizing clearance instead of minimizing travel distance




the example above is borrowed from “AI: A Modern Approach” by

Russell & P. Norvig

Planning via Skeletonization

- Voronoi diagrams
 - compute voronoi diagram ($O(n \log n)$, where n - # of invalid configurations)
 - add a shortest path segment from start to the nearest segment of voronoi diagram
 - add a shortest path segment from goal to the nearest segment of voronoi diagram
 - compute shortest path in the graph




the example above is borrowed from "AI: A Modern Approach" by

Russell & P. Norvig

Planning via Skeletonization

- Voronoi diagrams
 - advantages:
 - tends to stay away from obstacles
 - independent of the size of the environment
 - disadvantages:
 - can result in highly suboptimal paths




the example above is borrowed from “AI: A Modern Approach” by

Russell & P. Norvig

Planning via Skeletonization

- Voronoi diagrams
 - advantages:
 - tends to stay away from obstacles
 - independent of the size of the environment
 - disadvantages:
 - can result in highly suboptimal paths

In which environments?




the example above is borrowed from “AI: A Modern Approach” by

Russell & P. Norvig

Planning via Skeletonization

- Probabilistic roadmaps [Kavraki et al. '96]
 - construct a graph by:
 - randomly sampling valid configurations
 - adding edges in between the samples that are easy to connect with a straight line
 - add start and goal configurations to the graph with appropriate edges
 - compute shortest path in the graph




the example above is borrowed from “AI: A Modern Approach” by

Russell & P. Norvig

Planning via Skeletonization

- Probabilistic roadmaps [Kavraki et al. '96]
 - simple and highly effective (especially in $>2D$)
 - very popular
 - can result in suboptimal paths, no guarantees on suboptimality
 - difficulty with narrow passages




the example above is borrowed from “AI: A Modern Approach” by

Russell & P. Norvig

Planning via Skeletonization

- **Navmeshes**


- pick centers of triangles defining floor plan as graph vertices
- semi-manual but very popular in games
- can result in suboptimal paths, no guarantees on suboptimality



Planning via Skeletonization

- **Navmeshes**

- pick centers of triangles defining floor plan as graph vertices
- semi-manual but very popular in games
- can result in suboptimal paths, no guarantees on suboptimality



Other disadvantages?


Planning as Graph Search Problem

1. Construct a graph representing the planning problem
2. Search the graph for a (hopefully, close-to-optimal) path

The two steps above are often interleaved

Searching Graphs for a Least-cost Path


- Once a graph is constructed (from skeletonization or uniform cell decomposition or adaptive cell decomposition or lattice or whatever else), we need to search it for a least-cost path



A* Search

- Computes optimal g-values for relevant states


at any point of time:



A* Search

- Computes optimal g-values for relevant states

at any point of time:




one popular heuristic function – Euclidean distance

A* Search

- Heuristic function must be:
 - admissible: for every state s , $h(s) \leq c^*(s, s_{goal})$
 - consistent (satisfy triangle inequality):
 $h(s_{goal}, s_{goal}) = 0$ and for every $s \neq s_{goal}$, $h(s) \leq c(s, succ(s)) + h(succ(s))$
 - admissibility follows from consistency and often consistency follows from admissibility

minimal cost from s to s_{goal}



A* Search

- Computes optimal g-values for relevant states

Main function

$g(s_{start}) = 0$; all other g-values are infinite; $OPEN = \{s_{start}\}$;

ComputePath();

publish solution;

ComputePath function


while(s_{goal} is not expanded)

remove s with the smallest $[f(s) = g(s) + h(s)]$ from $OPEN$;

expand s ;

set of candidates for expansion

*for every expanded state
g(s) is optimal
(if heuristics are consistent)*



A* Search

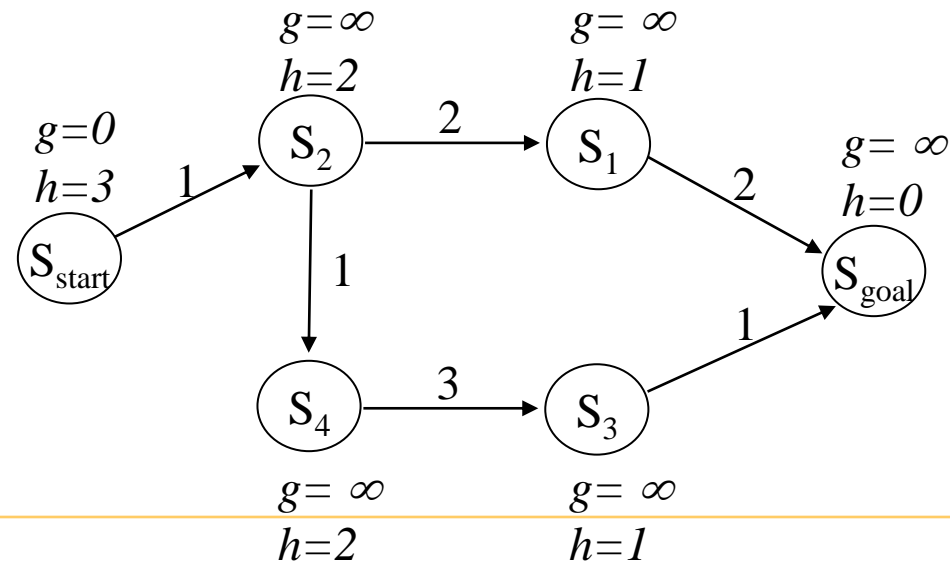
- Computes optimal g-values for relevant states

ComputePath function

while(s_{goal} is not expanded)

 remove s with the smallest $[f(s) = g(s) + h(s)]$ from *OPEN*;

 expand s ;



A* Search

- Computes optimal g-values for relevant states

ComputePath function

while(s_{goal} is not expanded)

remove s with the smallest $[f(s) = g(s) + h(s)]$ from *OPEN*;

insert s into *CLOSED*;

for every successor s' of s such that s' not in *CLOSED*

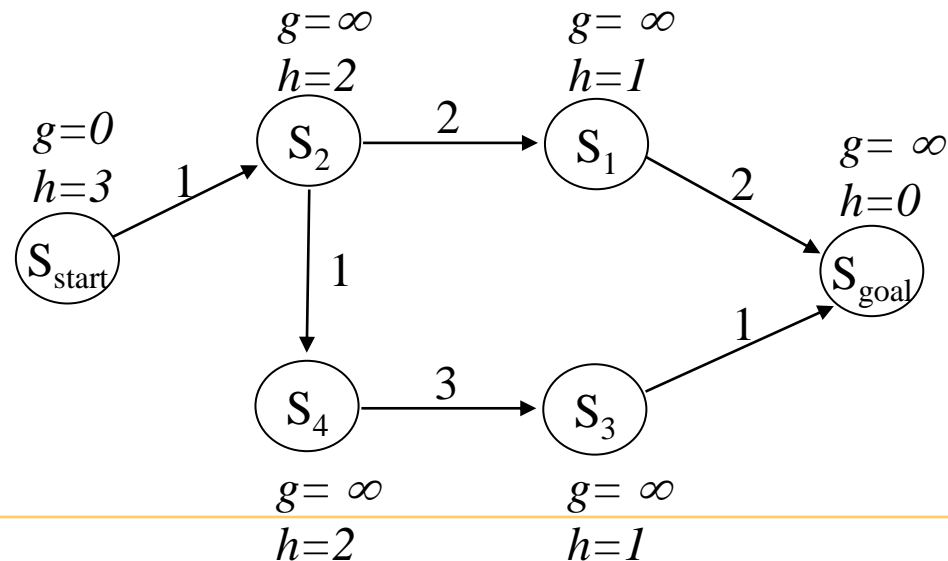
if $g(s') > g(s) + c(s, s')$

$g(s') = g(s) + c(s, s')$;

insert s' into *OPEN*;

tries to decrease $g(s')$ using the
found path from s_{start} to s

set of states that have already been expanded



A* Search

- Computes optimal g-values for relevant states

ComputePath function

while(s_{goal} is not expanded)

remove s with the smallest $[f(s) = g(s) + h(s)]$ from *OPEN*;

insert s into *CLOSED*;

for every successor s' of s such that s' not in *CLOSED*

if $g(s') > g(s) + c(s, s')$

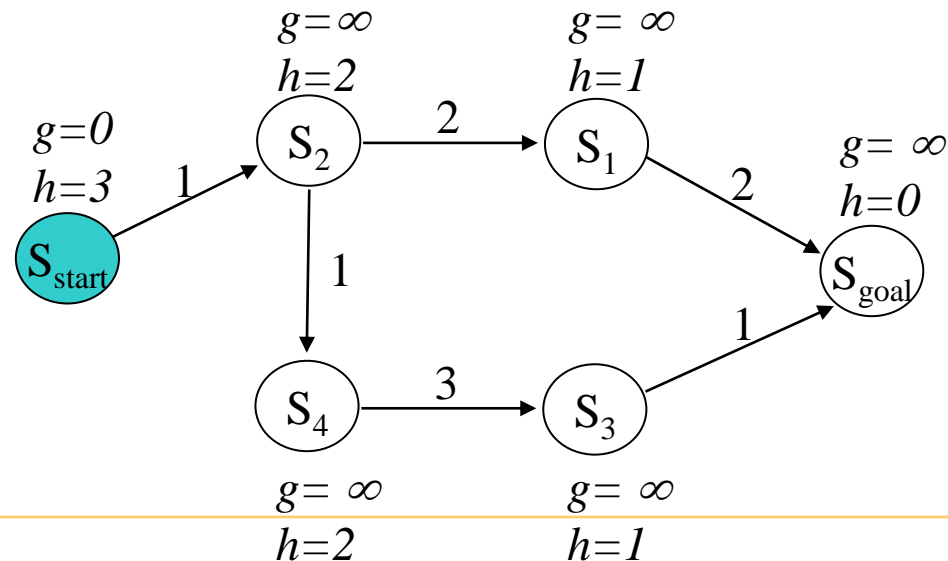
$g(s') = g(s) + c(s, s')$;

insert s' into *OPEN*;

CLOSED = {}

OPEN = { s_{start} }

next state to expand: s_{start}



A* Search

- Computes optimal g-values for relevant states

ComputePath function

while(s_{goal} is not expanded)

remove s with the smallest $[f(s) = g(s) + h(s)]$ from *OPEN*;

insert s into *CLOSED*;

for every successor s' of s such that s' not in *CLOSED*

if $g(s') > g(s) + c(s, s')$

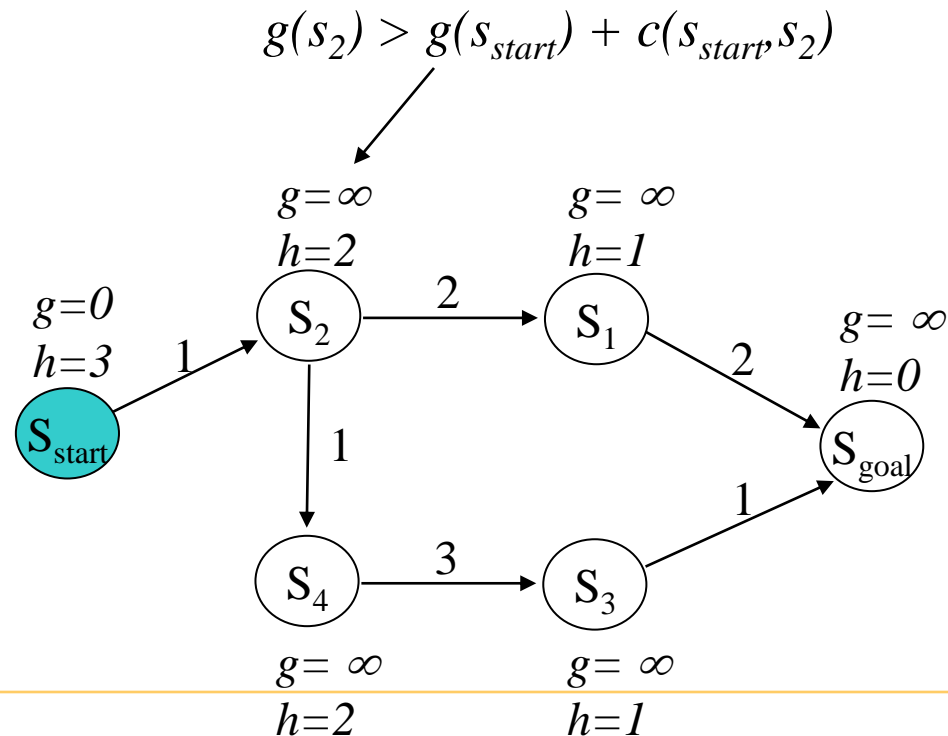
$g(s') = g(s) + c(s, s')$;

insert s' into *OPEN*;

CLOSED = {}

OPEN = { s_{start} }

next state to expand: s_{start}



A* Search

- Computes optimal g-values for relevant states

ComputePath function

while(s_{goal} is not expanded)

remove s with the smallest $[f(s) = g(s) + h(s)]$ from *OPEN*;

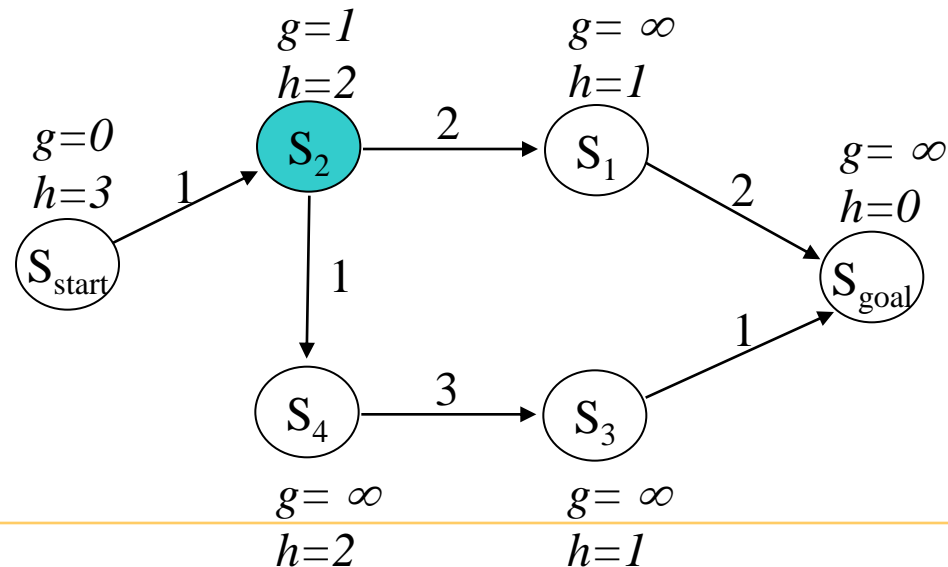
insert s into *CLOSED*;

for every successor s' of s such that s' not in *CLOSED*

if $g(s') > g(s) + c(s, s')$

$g(s') = g(s) + c(s, s')$;

insert s' into *OPEN*;



A* Search

- Computes optimal g-values for relevant states

ComputePath function

while(s_{goal} is not expanded)

remove s with the smallest $[f(s) = g(s) + h(s)]$ from *OPEN*;

insert s into *CLOSED*;

for every successor s' of s such that s' not in *CLOSED*

if $g(s') > g(s) + c(s, s')$


$g(s') = g(s) + c(s, s')$;

insert s' into *OPEN*;

CLOSED = $\{s_{start}\}$

OPEN = $\{s_2\}$

next state to expand: s_2



A* Search

- Computes optimal g-values for relevant states

ComputePath function

while(s_{goal} is not expanded)

remove s with the smallest $[f(s) = g(s) + h(s)]$ from *OPEN*;

insert s into *CLOSED*;

for every successor s' of s such that s' not in *CLOSED*

if $g(s') > g(s) + c(s, s')$


$g(s') = g(s) + c(s, s')$;

insert s' into *OPEN*;

CLOSED = $\{s_{start}, s_2\}$

OPEN = $\{s_1, s_4\}$

next state to expand: s_1



A* Search

- Computes optimal g-values for relevant states

ComputePath function

while(s_{goal} is not expanded)

remove s with the smallest $[f(s) = g(s) + h(s)]$ from *OPEN*;

insert s into *CLOSED*;

for every successor s' of s such that s' not in *CLOSED*

if $g(s') > g(s) + c(s, s')$


$g(s') = g(s) + c(s, s')$;

insert s' into *OPEN*;

CLOSED = $\{s_{start}, s_2, s_1\}$

OPEN = $\{s_4, s_{goal}\}$

next state to expand: s_4



A* Search

- Computes optimal g-values for relevant states

ComputePath function

while(s_{goal} is not expanded)

remove s with the smallest $[f(s) = g(s) + h(s)]$ from *OPEN*;

insert s into *CLOSED*;

for every successor s' of s such that s' not in *CLOSED*

if $g(s') > g(s) + c(s, s')$


$g(s') = g(s) + c(s, s')$;

insert s' into *OPEN*;

CLOSED = $\{s_{start}, s_2, s_1, s_4\}$

OPEN = $\{s_3, s_{goal}\}$

next state to expand: s_{goal}



A* Search

- Computes optimal g-values for relevant states

ComputePath function

while(s_{goal} is not expanded)

remove s with the smallest $[f(s) = g(s) + h(s)]$ from *OPEN*;

insert s into *CLOSED*;

for every successor s' of s such that s' not in *CLOSED*

if $g(s') > g(s) + c(s, s')$


$g(s') = g(s) + c(s, s')$;

insert s' into *OPEN*;

CLOSED = $\{s_{start}, s_2, s_1, s_4, s_{goal}\}$

OPEN = $\{s_3\}$

done



A* Search

- Computes optimal g-values for relevant states

ComputePath function

while(s_{goal} is not expanded)

remove s with the smallest $[f(s) = g(s) + h(s)]$ from *OPEN*;


insert s into *CLOSED*;

for every successor s' of s such that s' not in *CLOSED*

if $g(s') > g(s) + c(s, s')$

$g(s') = g(s) + c(s, s')$;

insert s' into *OPEN*;



for every expanded state $g(s)$ is optimal
for every other state $g(s)$ is an upper bound
we can now compute a least-cost path

A* Search

- Computes optimal g-values for relevant states

ComputePath function

while(s_{goal} is not expanded)

remove s with the smallest $[f(s) = g(s) + h(s)]$ from *OPEN*;


insert s into *CLOSED*;

for every successor s' of s such that s' not in *CLOSED*

if $g(s') > g(s) + c(s, s')$

$g(s') = g(s) + c(s, s')$;

insert s' into *OPEN*;



A* Search

- Computes optimal g-values for relevant states

ComputePath function

while(s_{goal} is not expanded)

remove s with the smallest $[f(s) = g(s) + h(s)]$ from *OPEN*;

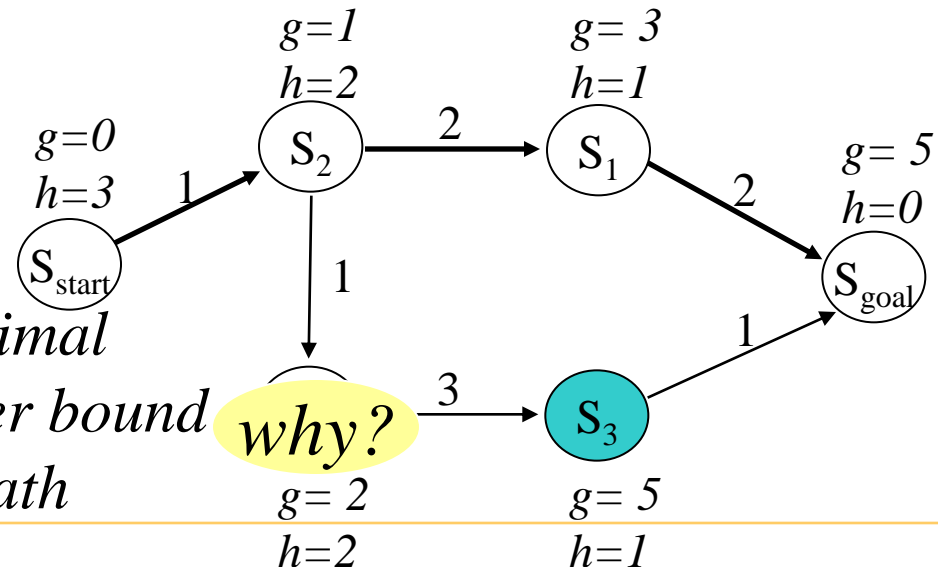
insert s into *CLOSED*;

for every successor s' of s such that s' not in *CLOSED*

if $g(s') > g(s) + c(s, s')$

$g(s') = g(s) + c(s, s')$;

insert s' into *OPEN*;




for every expanded state $g(s)$ is optimal

for every other state $g(s)$ is an upper bound

we can now compute a least-cost path

A* Search

- Is guaranteed to return an optimal path (in fact, for every expanded state) – optimal in terms of the solution
- Performs provably minimal number of state expansions required to guarantee optimality – optimal in terms of the computations



Effect of the Heuristic Function

- A* Search: expands states in the order of $f = g + h$ values

ComputePath function

while(s_{goal} is not expanded)

 remove s with the smallest $[f(s) = g(s) + h(s)]$ from *OPEN*;

 insert s into *CLOSED*;

 for every successor s' of s such that s' not in *CLOSED*

 if $g(s') > g(s) + c(s, s')$

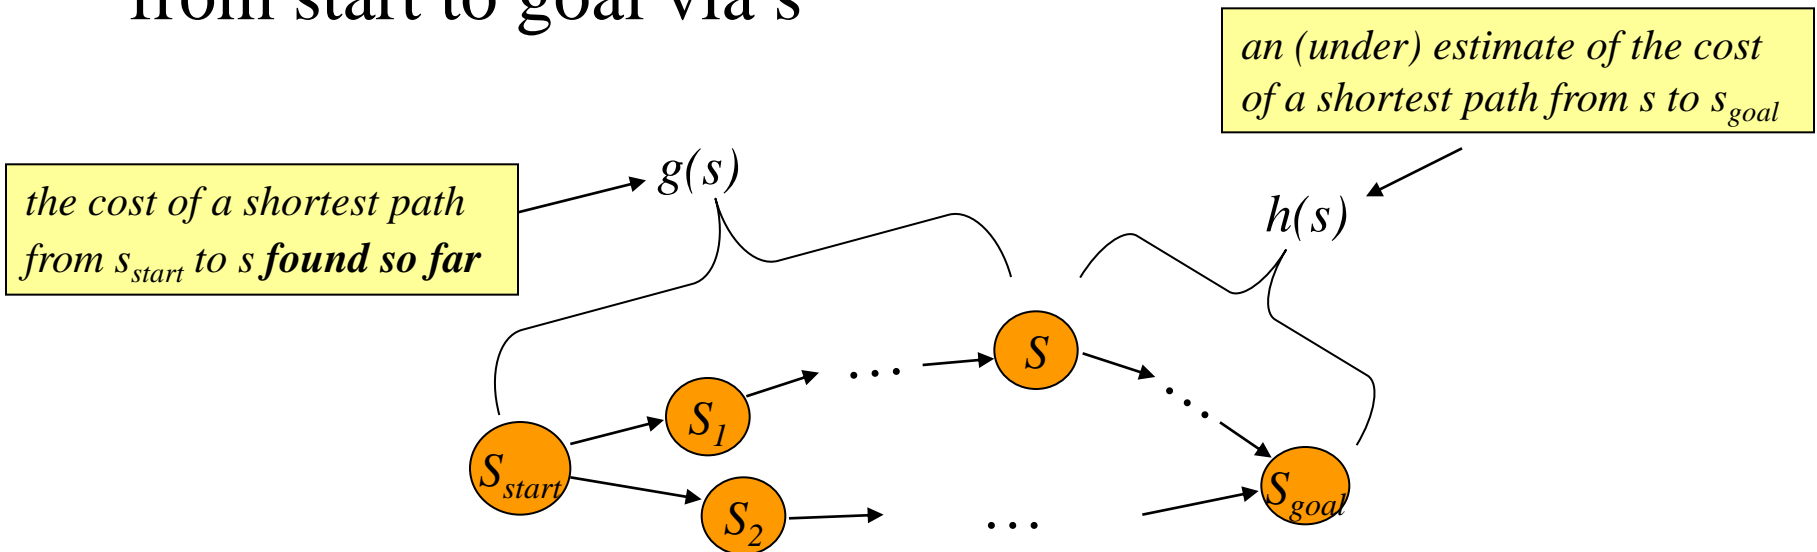
$g(s') = g(s) + c(s, s')$;

 insert s' into *OPEN*;

} *expansion of s*


Effect of the Heuristic Function

- A* Search: expands states in the order of $f = g + h$ values
- Dijkstra's: expands states in the order of $f = g$ values (pretty much)
- Intuitively: $f(s)$ – estimate of the cost of a least cost path from start to goal via s



Effect of the Heuristic Function


- **A* Search:** expands states in the order of $f = g + h$ values
- **Dijkstra's:** expands states in the order of $f = g$ values (pretty much)
- **Weighted A*:** expands states in the order of $f = g + \varepsilon h$ values, $\varepsilon > 1$ = bias towards states that are closer to goal



Effect of the Heuristic Function

- Dijkstra's: expands states in the order of $f = g$ values


What are the states expanded?



Effect of the Heuristic Function

- A* Search: expands states in the order of $f = g + h$ values


What are the states expanded?



Effect of the Heuristic Function

- A* Search: expands states in the order of $f = g + h$ values


for large problems this results in A being slow*



Effect of the Heuristic Function

- Weighted A* Search: expands states in the order of $f = g + \epsilon h$ values, $\epsilon > 1$ = bias towards states that are closer to goal

*what states are expanded?
– research question*



Effect of the Heuristic Function

- Weighted A* Search:


- trades off optimality for speed

- ϵ -suboptimal:

$$\text{cost}(\text{solution}) \leq \epsilon \cdot \text{cost}(\text{optimal solution})$$


- in many domains, it has been shown to be orders of magnitude faster than A*

A: $\epsilon = 1.0$*



*20 expansions
solution=10 moves*

Weighted A: $\epsilon = 2.5$*



*13 expansions
solution=11 moves*